

A case study in HR analytics

Shubham

Steps for text mining project

- Problem Defination and Specific Goal
- Identify Text to be collected
- Text Organization
- Feature Extraction
- Analysis
- Reach an insight

Problem :- Which company has better work life balance? Which has better perceived pay according to online reviews? We learn something about how employees review both Amazon and Google.

Employee reviews can come from various sources. Forbes and others publish articles about the “best places to work”, which may mention Amazon and Google. Another source of information might be anonymous online reviews from websites like Indeed, Glassdoor or CareerBliss.

Here, we'll focus on a collection of anonymous online reviews of amazon and google.

```
library(readr)
amazon <- read_csv("~/500_amzn.csv")
google <- read_csv("~/500_goog.csv")
```

```
str(amazon)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   500 obs. of  4 variables:
## $ pg_num: int  50 50 50 50 50 50 50 50 50 50 ...
## $ url : chr  "https://www.glassdoor.com/Reviews/Amazon-com-Reviews-E6036_P50.htm" "https://www.glassdoor.com/Reviews/Google-Reviews-E9079_P1.htm"
## $ pros : chr  "You're surrounded by smart people and the projects are interesting, if a little daunting" "If you're a software engineer, you're among the kings of the hill at Google. It's a great place to work"
## $ cons : chr  "Internal tools proliferation has created a mess for trying to get to basic information" "The work is very repetitive and boring"
## - attr(*, "spec")=List of 2
## ..$ cols :List of 4
## .. ..$ pg_num: list()
## .. .. ..- attr(*, "class")= chr  "collector_integer" "collector"
## .. ..$ url : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## .. ..$ pros : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## .. ..$ cons : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## ..$ default: list()
## .. ..- attr(*, "class")= chr  "collector_guess" "collector"
## ..- attr(*, "class")= chr  "col_spec"
```

```
amazon_pros <- amazon$pros
amazon_cons <- amazon$cons
```

```
str(google)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   501 obs. of  4 variables:
## $ pg_num: int  1 1 1 1 1 1 1 1 1 1 ...
## $ url : chr  "https://www.glassdoor.com/Reviews/Google-Reviews-E9079_P1.htm" "https://www.glassdoor.com/Reviews/Google-Reviews-E9079_P1.htm"
## $ pros : chr  "* If you're a software engineer, you're among the kings of the hill at Google. It's a great place to work" "If you're a software engineer, you're among the kings of the hill at Google. It's a great place to work"
```

```
## $ cons : chr  "* It *is* becoming larger, and with it comes growing pains: bureaucracy, slow to re
## - attr(*, "spec")=List of 2
## ..$ cols :List of 4
## .. ..$ pg_num: list()
## .. .. ..- attr(*, "class")= chr  "collector_integer" "collector"
## .. ..$ url : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## .. ..$ pros : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## .. ..$ cons : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## ..$ default: list()
## .. ..- attr(*, "class")= chr  "collector_guess" "collector"
## ..- attr(*, "class")= chr "col_spec"
```

```
google_pros <- google$pros
google_cons <- google$cons
```

Text organization Now that we have selected the exact text sources, we are now ready to clean them up. We'll be using the two functions `qdap_clean()`, which applies a series of `qdap` functions to a text vector, and `tm_clean()`, which applies a series of `tm` functions to a corpus object.

```
library(qdap)
qdap_clean <- function(x){
  x<- na.omit(x)
  x<- replace_abbreviation(x)
  x<- replace_contraction(x)
  x<- replace_number(x)
  x<- replace_ordinal(x)
  x<- replace_symbol(x)
  x<-tolower(x)
  return(x)
}
```

```
library(tm)
```

```
tm_clean <- function(x){
  x<-tm_map(x,removePunctuation)
  x<-tm_map(x,stripWhitespace)
  x<-tm_map(x,removeWords,c(stopwords("en"),"Amazon","Google","Company"))
  return(x)
}
```

Applying `qdap_clean()` to amazon and google

```
amazon_pros <- qdap_clean(amazon_pros)
amazon_cons <- qdap_clean(amazon_cons)

google_pros <- qdap_clean(google_pros)
google_cons <- qdap_clean(google_cons)
```

Next step is to convert this vector containing the text data to a corpus. Corpus is a collection of documents, but it's also important to know that in the `tm` domain, R recognizes it as a data type.

There are two kinds of the corpus data type, the permanent corpus, `PCorpus`, and the volatile corpus, `VCorpus`. In essence, the difference between the two has to do with how the collection of documents is stored in your computer. We will use the volatile corpus, which is held in computer's RAM rather than saved to disk,

just to be more memory efficient.

To make a volatile corpus, R needs to interpret each element in our vector of text, `amazon_pros`, as a document. And the `tm` package provides what are called Source functions to do just that! We'll use a Source function called `VectorSource()` because our text data is contained in a vector. The output of this function is called a Source.

```
amazon_p_corp <- VCorpus(VectorSource(amazon_pros))
amazon_c_corp <- VCorpus(VectorSource(amazon_cons))

google_p_corp <- VCorpus(VectorSource(google_pros))
google_c_corp <- VCorpus(VectorSource(google_cons))
```

Now using `tm_clean` to clean data

```
amazon_pros_corp <- tm_clean(amazon_p_corp)
amazon_cons_corp <- tm_clean(amazon_c_corp)

google_pros_corp <- tm_clean(google_p_corp)
google_cons_corp <- tm_clean(google_c_corp)
```

Steps 4 & 5: Feature extraction & analysis

Since `amzn_pros_corp`, `amzn_cons_corp`, `goog_pros_corp` and `goog_cons_corp` have all been preprocessed, so now we can extract the features we want to examine. Since we are using the bag of words approach, we decide to create a bigram `TermDocumentMatrix` for Amazon's positive reviews corpus, `amzn_pros_corp`. From this, we can quickly create a `wordcloud()` to understand what phrases people positively associate with working at Amazon.

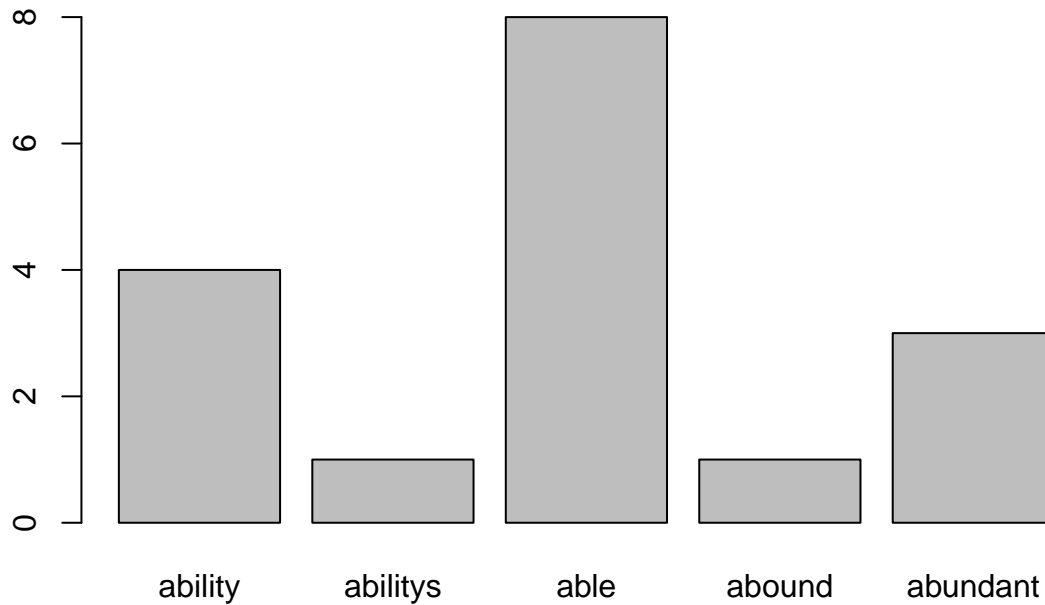
The function below uses `RWeka` to tokenize two terms.

```
library(RWeka)
tokenizer <- function(x)
  NGramTokenizer(x, Weka_control(min = 2, max = 2))
```

Feature extraction & analysis: `amazon_cons`

```
amazon_p_tdm <- TermDocumentMatrix(amazon_pros_corp)
amazon_p_tdm_m <- as.matrix(amazon_p_tdm)
amazon_p_freq <- rowSums(amazon_p_tdm_m)
amazon_p_f.sort <- sort(amazon_p_freq, decreasing = TRUE)

barplot(amazon_p_freq[1:5])
```



```
library(wordcloud)
amazon_p_tdm <- TermDocumentMatrix(amazon_pros_corp, control = list(tokenize=tokenizer))
amazon_p_tdm_m <- as.matrix(amazon_p_tdm)
amazon_p_freq <- rowSums(amazon_p_tdm_m)
amazon_p_f.sort <- sort(amazon_p_freq, decreasing = TRUE)
p_df <- data.frame(term=names(amazon_p_f.sort), num=amazon_p_f.sort)

wordcloud(p_df$term, p_df$num, max.words=100, color="red")
```

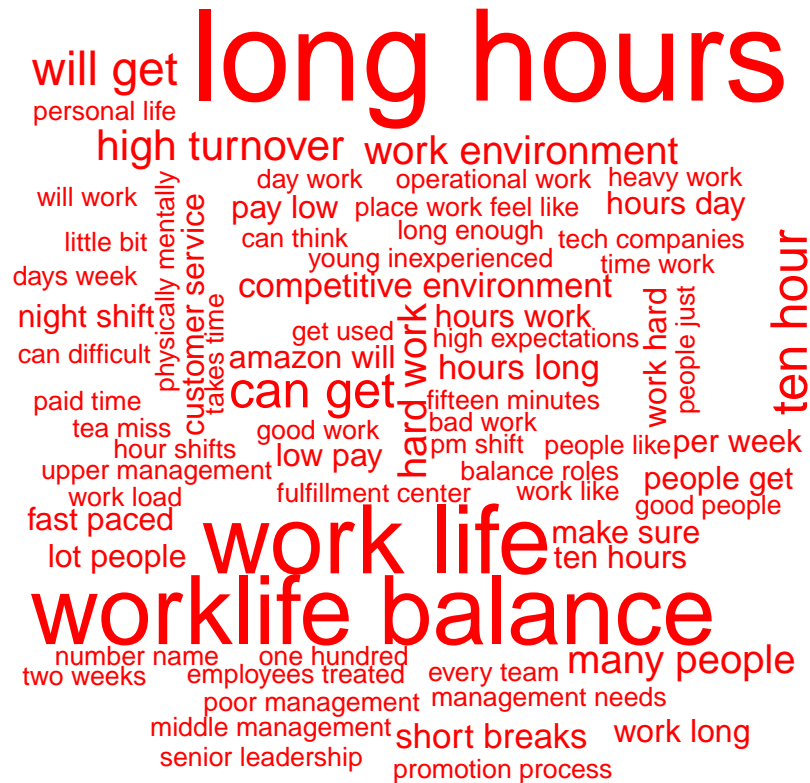


Feature extraction & analysis: ama-

zon_cons

```
amazon_c_tdm <- TermDocumentMatrix(amazon_cons_corp, control=list(tokenize=tokenizer))
amazon_c_tdm_m <- as.matrix(amazon_c_tdm)
amazon_c_freq <- rowSums(amazon_c_tdm_m)
amazon_c_f.sort <- sort(amazon_c_freq, decreasing = TRUE)
c_df <- data.frame(term=names(amazon_c_f.sort), num=amazon_c_f.sort)

wordcloud(c_df$term, c_df$num, max.words=100, color="red")
```



amazon_cons dendrogram

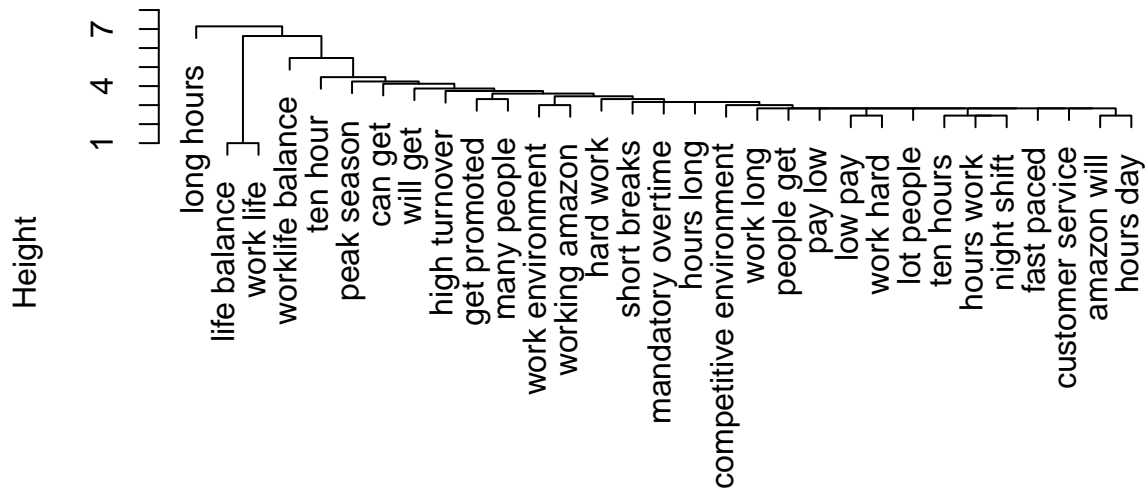
It seems there is a strong indication of long working hours and poor work-life balance in the reviews. As a simple clustering technique, we'll decide to perform a hierarchical cluster and create a dendrogram to see how connected these phrases are.

```
amazon_c_tdm <- TermDocumentMatrix(amazon_cons_corp, control = list(tokenize=tokenizer))
amazon_c_tdm <- removeSparseTerms(amazon_c_tdm, 0.993)

amazon_c_hclust <- hclust(dist(amazon_c_tdm, method="euclidean"), method="complete")

plot(amazon_c_hclust)
```

Cluster Dendrogram



```
dist(amazon_c_tdm, method = "euclidean")
hclust (*, "complete")
```

Word association Switching back to positive comments, we'll decide to examine top phrases that appeared in the word clouds. We'll now hope to find associated terms using the findAssocs() function from tm package.

```
amazon_p_tdm <- TermDocumentMatrix(amazon_pros_corp, control=list(tokenize=tokenizer))
amazon_p_m <- as.matrix(amazon_p_tdm)
amazon_p_freq <- rowSums(amazon_p_m)
token_frequency <- sort(amazon_p_freq, decreasing = TRUE)
token_frequency[1:5]
```

```
##      good pay great benefits    smart people    place work    fast paced
##           25             24             20             17             16
```

```
findAssocs(amazon_p_tdm, "fast paced", 0.2)
```

```
## $`fast paced`
##      paced environment    environments ever    learn fast
##           0.49             0.35             0.35
##      paced friendly      paced work          able excel
##           0.35             0.35             0.25
##      activity ample      advance one         also well
##           0.25             0.25             0.25
##      amazon fast         amazon noting       amazon one
##           0.25             0.25             0.25
##      amount time        ample opportunity    assistance ninety
##           0.25             0.25             0.25
##      benefits including    break computer    call activity
##           0.25             0.25             0.25
##      can choose           catchy company      center things
##           0.25             0.25             0.25
##      challenging expect    cheers opportunity    choose success
```

##	0.25	0.25	0.25
##	combined encouragement	company cheers	competitive environments
##	0.25	0.25	0.25
##	computer room	cool things	deliver results
##	0.25	0.25	0.25
##	dock makes	driven deliver	easy learn
##	0.25	0.25	0.25
##	emphasis shipping	encouragement innovation	environment benefits
##	0.25	0.25	0.25
##	environment catchy	environment center	environment fast
##	0.25	0.25	0.25
##	environment help	environment smart	everchanging fast
##	0.25	0.25	0.25
##	ever known	ever witnessed	everyones preferences
##	0.25	0.25	0.25
##	excel advance	excel everchanging	exciting environment
##	0.25	0.25	0.25
##	expect learn	extremely fast	facility top
##	0.25	0.25	0.25
##	fail successful	fantastic able	fired part
##	0.25	0.25	0.25
##	five percent	freindly place	friendly atmosphere
##	0.25	0.25	0.25
##	friendly management	full medical	get fired
##	0.25	0.25	0.25
##	go extremely	great plenty	great teamwork
##	0.25	0.25	0.25
##	happening technology	hassle benefits	help get
##	0.25	0.25	0.25
##	help workers	high quality	high volume
##	0.25	0.25	0.25
##	including full	innovation owning	job requirements
##	0.25	0.25	0.25
##	leader can	line break	lot responsibility
##	0.25	0.25	0.25
##	maintaining high	makes time	management nice
##	0.25	0.25	0.25
##	nice facility	ninety five	noting short
##	0.25	0.25	0.25
##	offers opportunity	one competitive	one fast
##	0.25	0.25	0.25
##	opportunity overtime	opportunity yell	ownership fast
##	0.25	0.25	0.25
##	owning work	paced emphasis	paced exciting
##	0.25	0.25	0.25
##	paced high	paced never	paced rewarding
##	0.25	0.25	0.25
##	paced ship	paced software	paid upfront
##	0.25	0.25	0.25
##	people focused	percent paid	plenty shifts
##	0.25	0.25	0.25
##	position fast	possible still	preferences fast
##	0.25	0.25	0.25
##	products quickly	quality bar	quickly possible

##	0.25	0.25	0.25
##	readily available	requirements easy	responsibility ownership
##	0.25	0.25	0.25
##	results great	results team	rewarding people
##	0.25	0.25	0.25
##	shifts everyones	ship dock	shipping products
##	0.25	0.25	0.25
##	short amount	short fantastic	smart coworkers
##	0.25	0.25	0.25
##	still maintaining	success fail	successful also
##	0.25	0.25	0.25
##	team driven	technology today	things happening
##	0.25	0.25	0.25
##	things lot	time fast	time go
##	0.25	0.25	0.25
##	top line	upfront experience	vision well
##	0.25	0.25	0.25
##	volume call	well rewarded	well tuition
##	0.25	0.25	0.25
##	witnessed combined	work can	work cool
##	0.25	0.25	0.25
##	work environments	workers readily	work fast
##	0.25	0.25	0.25
##	work job	yell leader	
##	0.25	0.25	

We decide to create a `comparison.cloud()` of Google's positive and negative reviews for comparison to Amazon. This will give you a quick understanding of top terms.

```
all_google_pros <- paste(goOGLE$pros,collapse="")
all_google_cons <- paste(goOGLE$cons,collapse = "")

all_google <- c(all_google_pros,all_google_cons)
all_google_qdap <- qdap_clean(all_google)
all_google_vs <- VectorSource(all_google_qdap)
all_google_vc <- VCorpus(all_google_vs)
all_google_clean<- tm_clean(all_google_vc)
all_google_tdm <- TermDocumentMatrix(all_google_clean)
colnames(all_google_tdm) <- c("Google Pros","Google Cons")
all_google_tdm_m <- as.matrix(all_google_tdm)

comparison.cloud(all_google_tdm_m,colors = c("orange","blue"),max.words = 50)
```




Google Cons

Amazon's positive reviews appear to mention bigrams such as "good benefits", while its negative reviews focus on bigrams such as "work-life balance" issues.

In contrast to, Google's positive reviews mention "perks", "smart people", "great food", and "fun culture", among other things. Google's negative reviews discuss "politics", "getting big", "bureaucracy", and "middle management".

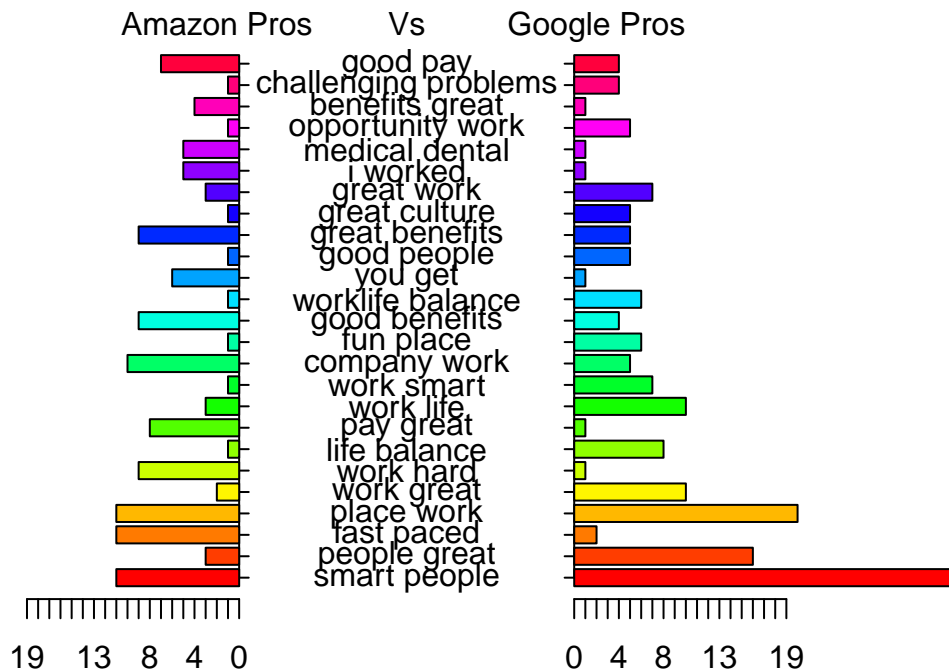
Now we'll make a pyramid plot lining up positive reviews for Amazon and Google so you can adequately see the differences between any shared bigrams.

```
amazon_pro <- paste(amazon$pros,collapse = "")
google_pro <- paste(google$pros,collapse = "")
all_pro <- c(amazon_pro,google_pro)
all_pro_qdap <- qdap_clean(all_pro)
all_pro_vs <- VectorSource(all_pro)
all_pro_vc <- VCorpus(all_pro_vs)
all_pro_corp <- tm_clean(all_pro_vc)

tdm.bigram = TermDocumentMatrix(all_pro_corp,control = list(tokenize =tokenizer))
colnames(tdm.bigram) <- c("Amazon","Google")
tdm.bigram <- as.matrix(tdm.bigram)
common_words<- subset(tdm.bigram,tdm.bigram[,1] > 0 & tdm.bigram[,2] > 0 )
difference <- abs(common_words[, 1] - common_words[,2])
common_words <- cbind(common_words,difference)
common_words <- common_words[order(common_words[,3],decreasing = TRUE),]
top25_df <- data.frame(x=common_words[1:25,1],y=common_words[1:25,2],labels=rownames(common_words[1:25,]))

library(plotrix)
pyramid.plot(top25_df$x,top25_df$y,labels=top25_df$labels,gap=15,top.labels=c("Amazon Pros","Vs","Google Cons"))
```

Words in common



```
## [1] 5.1 4.1 4.1 2.1
```

Amazon employees discussed “work-life balance” as a positive. In both organizations, people mentioned “culture” and “smart people”, so there are some similar positive aspects between the two companies.

You now decide to turn your attention to negative reviews and make the same visuals.

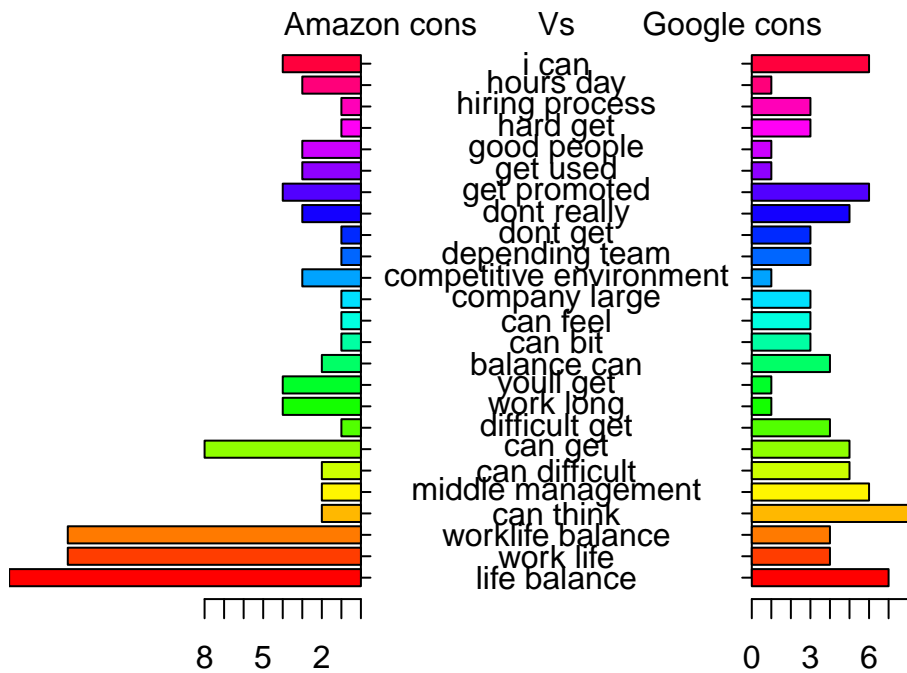
```
amazon_cons <- paste(amazon$cons, collapse = "")
google_cons <- paste(google$cons, collapse = "")
all_cons <- c(amazon_cons, google_cons)
all_cons_qdap <- qdap_clean(all_cons)
all_cons_vs <- VectorSource(all_cons)
all_cons_vc <- VCorpus(all_cons_vs)
all_cons_corp <- tm_clean(all_cons_vc)

tdm.cons_bigram = TermDocumentMatrix(all_cons_corp, control=list(tokenize = tokenizer))

colnames(tdm.cons_bigram) <- c("Amazon", "Google")
tdm.cons_bigram <- as.matrix(tdm.cons_bigram)
common_words <- subset(tdm.cons_bigram, tdm.cons_bigram[,1] > 0 & tdm.cons_bigram[,2] > 0)
difference <- abs(common_words[, 1] - common_words[, 2])
common_words <- cbind(common_words, difference)
common_words <- common_words[order(common_words[, 3], decreasing = TRUE), ]
top25_df <- data.frame(x=common_words[1:25, 1], y=common_words[1:25, 2], labels=rownames(common_words[1:25, ]))

library(plotrix)
pyramid.plot(top25_df$x, top25_df$y, labels=top25_df$labels, gap=10, top.labels=c("Amazon cons", "Vs", "Google cons"))
```

Words in common



```
## [1] 5.1 4.1 4.1 2.1
```

We'll use Commonality cloud to show common between Aamazon and google with Unigram, Bigram and Trigram tokenizer to identify more insights.

Unigram

```
tdm.unigram <- TermDocumentMatrix(all_pro_corp)
colnames(tdm.unigram) <- c("Amazon", "Google")
tdm.unigram <- as.matrix(tdm.unigram)

commonality.cloud(tdm.unigram, colors=c("red", "yellow"), max.words = 100)
```



Bigram

```
BigramTokenizer <- function(x) NGramTokenizer(x, Weka_control(min = 2, max = 2))
tdm.bigram <- TermDocumentMatrix(all_pro_corp, control = list(tokenize=BigramTokenizer))
colnames(tdm.bigram) <- c("Amazon", "Google")
tdm.bigram <- as.matrix(tdm.bigram)

commonality.cloud(tdm.bigram, colors=c("red", "yellow"), max.words = 100)
```



Trigram

```
TrigramTokenizer <- function(x) NGramTokenizer(x, Weka_control(min = 3, max = 3))
tdm.trigram <- TermDocumentMatrix(all_pro_corp, control = list(tokenize=TrigramTokenizer))
colnames(tdm.trigram) <- c("Amazon", "Google")
tdm.trigram <- as.matrix(tdm.trigram)

commonality.cloud(tdm.trigram, colors=c("red", "yellow"), max.words = 100)
```

people great work
company great people

work a lot
good work life

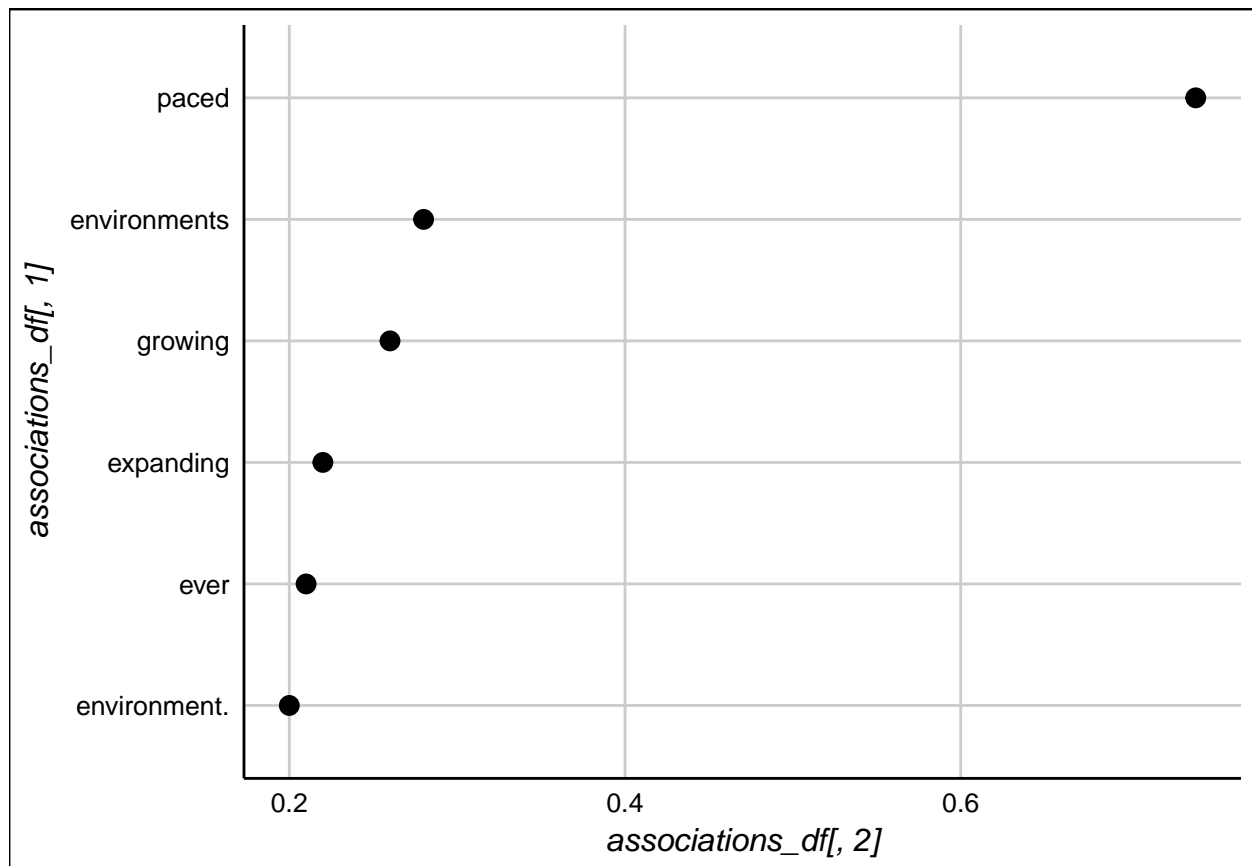
really enjoyed working
work life balancegreat

Plotting amazon and google association

```
library(ggthemes)
library(ggplot2)

amazon_tdm <- TermDocumentMatrix(amazon_p_corp)
associations <- findAssocs(amazon_tdm,"fast",0.2)
associations_df <- list_vect2df(associations)[,2:3]

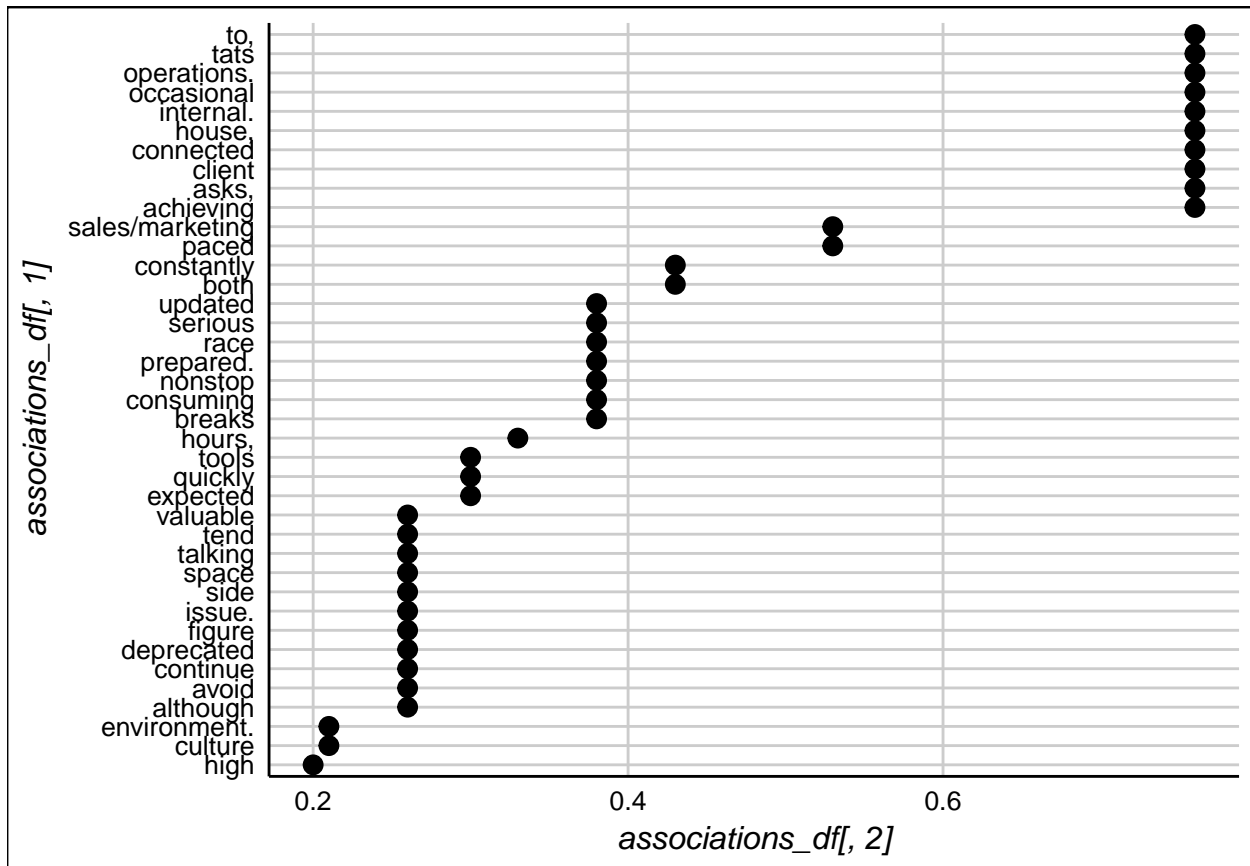
ggplot(associations_df,aes(y=associations_df[,1]))+
  geom_point(aes(x=associations_df[,2]),
             data=associations_df,size=3)+
  theme_gdocs()
```



```
library(ggthemes)
library(ggplot2)

google_tdm <- TermDocumentMatrix(google_c_corp)
associations <- findAssocs(google_tdm, "fast", 0.2)
associations_df <- list_vect2df(associations)[,2:3]

ggplot(associations_df, aes(y=associations_df[,1]))+
  geom_point(aes(x=associations_df[,2]),
             data=associations_df, size=3)+
  theme_gdocs()
```



Conclusion Google have a better work-life balance according to current employee reviews.

```
findAssocs(amazon_p_tdm, "fast paced", 0.2)[[1]][1:15]
```

```
## paced environment environments ever      learn fast      paced friendly
##           0.49           0.35           0.35           0.35
##      paced work      able excel      activity ample      advance one
##           0.35           0.25           0.25           0.25
##      also well      amazon fast      amazon noting      amazon one
##           0.25           0.25           0.25           0.25
##      amount time ample opportunity assistance ninety
##           0.25           0.25           0.25
```

We Identified candidates that view an intense workload as an opportunity to learn fast and give them ample opportunity.